

1. Write a MapReduce program to count the number of occurrences of each word and provide output as follows:

Output

Word Word Count

a 1 (As the word 'a' occurred only once)

this 2 (As the word 'this' occurred twice)

```
package cdac_hadoop;

import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class WCmapper extends Mapper<LongWritable, Text, Text, IntWritable>{
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(LongWritable key, Text value, Context context) throws
IOException, InterruptedException{
        StringTokenizer itr = new StringTokenizer(value.toString());

        while (itr.hasMoreTokens()){
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}

package cdac_hadoop;

import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class WCReducer extends Reducer<Text, IntWritable, Text, IntWritable>{
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws
IOException, InterruptedException
    {
        int sum = 0;
        for (IntWritable val: values){
            sum = sum + val.get();
        }
    }
}
```

```

        result.set(sum);
        context.write(key, result);
    }

}

package cdac_hadoop;

import org.apache.hadoop.conf.Configuration;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
public class WordCount {

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word Count");

        job.setJarByClass(WordCount.class);
        job.setMapperClass(WCmapper.class);
        job.setReducerClass(WCReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        System.exit(job.waitForCompletion(true)? 0 : 1);
    }

}

```

Export java code as jar file.

Put the file in hdfs system

Command :- `hdfs dfs -put computer.txt /user/osboxes/input/`

Then use the jar file to run perform mapreduce program on above loaded text file.

Hadoop jar `WC.jar /user/osboxes/input/computer.txt /user/osboxes/output/`

Output:-

```
osboxes@osboxes: ~/Downloads
osboxes@osboxes:~$ hdfs -cat /user/osboxes/output/part:r-00000
2021-01-13 11:03:34,750 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
"complete" 1
"full" 1
(CPU) 1
(IO) 1
(MOS) 1
(as) 1
(e.g.) 1
(keyboards, 1
(main 1
(monitor 1
1940s 1
1950s 1
1970s 1
2000s-era 1
20th 2
21st 1
A 2
Computers 1
Conventionally, 1
Digital 1
Early 2
II, 1
Industrial 1
Internet 1
MOORE 1
Modern 1
Moore's 1
More 1
Peripheral 2
Revolution 1
Revolution, 1
Since 1
The 5
These 1
This 2
War 1
World 1
a 10
abacus 1
ability 1
aided 1
allow 1
along 1
also 1
an 2
```

2. Write a MapReduce program that reads the alphabetst ext file and counts the occurrences of words of each size. The output should appear as follows:

Sample Output

Word Size Word Count

1 1 (As the word of size 1 is: a)

2 4 (As the words of size 2 are: is, of, of, in)

3 3 (As the words of size 3 are: the, and, the)

4 6 (As the words of size 4 are: this, word, size, that, size)

```
package Wordlength;
```

```
import java.io.IOException;
```

```
import org.apache.hadoop.mapreduce.Mapper;
```

```
import org.apache.hadoop.mapreduce.Reducer;
```

```
import org.apache.hadoop.mapreduce.Job;
```

```
import org.apache.hadoop.io.LongWritable;
```

```
import org.apache.hadoop.io.IntWritable;
```

```
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.conf.Configuration;
```

```
import org.apache.hadoop.fs.Path;
```

```
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
```

```
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```

public class AlphabetWordCount {

    public static class AlphabetWordCountMapper extends
    Mapper<LongWritable,Text,IntWritable,IntWritable>{

        // here we are declaring a final static local variable of static type

        public static final IntWritable ONE = new IntWritable(1);

        @Override

        public void map(LongWritable key, Text value, Context cont) throws
        IOException,InterruptedException{

            String line = value.toString();

            for (String word : line.split(" ")) {

                System.out.println("This is a Alphabet count problem");

                cont.write(new IntWritable(word.length()), ONE);

            }

        }

    }

    public static class AlphabetWordCountReducer extends Reducer<IntWritable, IntWritable,
    IntWritable, IntWritable>{

        @Override

        public void reduce(IntWritable key, Iterable<IntWritable> value, Context cont) throws
        IOException,InterruptedException{

            int count = 0;

            for (IntWritable values : value) {

                count += values.get();

            }

            cont.write(key, new IntWritable(count));

        }

    }

}

```

```

public static void main(String[] args) throws
ClassNotFoundException,IOException,InterruptedException{

    // TODO Auto-generated method stub

    Configuration conf = new Configuration();

    Job job = Job.getInstance(conf, "WordCount");

    job.setMapperClass(AlphabetWordCountMapper.class);

    job.setReducerClass(AlphabetWordCountReducer.class);

    job.setCombinerClass(AlphabetWordCountReducer.class);

    job.setMapOutputKeyClass(IntWritable.class);

    job.setMapOutputValueClass(IntWritable.class);

    job.setOutputKeyClass(IntWritable.class);

    job.setOutputValueClass(IntWritable.class);

    job.setJarByClass(AlphabetWordCount.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));

    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    System.exit(job.waitForCompletion(true) ? 0 : 1);

}
}

```

Export java code as jar file.

Put the file in hdfs system

Command :- `hdfs dfs -put computer.txt /user/osboxes/input/`

Then use the jar file to run perform mapreduce program on above loaded text file.

Hadoop jar `Wordlengh.jar /user/osboxes/input/computer.txt /user/osboxes/output1/`

Output:-

3. Write a MapReduce program to process a given patent dataset with patent records. Each patent has sub-patent id's associated with it. You need to calculate the number of sub-patents associated with each patent.

package patent;

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

/**
 * @author osboxes
 */
public class patentmapper extends Mapper<LongWritable,Text,Text,IntWritable> {
    private final static IntWritable one = new IntWritable(1);
    private Text key1 = new Text();
    private Text value1 = new Text();

    public void map(LongWritable key,Text value,Context context) throws
    IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());

        while (itr.hasMoreTokens()) {
            String patent = itr.nextToken();
            key1.set(patent);
```

```

        String subpatent = itr.nextTok
        value1.set(subpatent);
        context.write(key1, one);
    }
}

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class pat {

    public static void main(String[] args) throws Exception {
        // TODO Auto-generated method stub

        Configuration conf = new Configuration();

        Job job = Job.getInstance(conf, "patent");

        job.setJarByClass(pat.class);
        job.setMapperClass(patentmapper.class);
        job.setReducerClass(patred.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class pat {

    public static void main(String[] args) throws Exception {
        // TODO Auto-generated method stub

```

```

Configuration conf = new Configuration();

Job job = Job.getInstance(conf, "patent");

job.setJarByClass(pat.class);
job.setMapperClass(patentmapper.class);
job.setReducerClass(patred.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);

FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));

System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

```

Export java code as jar file.

Put the file in hdfs system

Command :- `hdfs dfs -put patent /user/osboxes/input/`

Then use the jar file to run perform mapreduce program on above loaded text file.

Hadoop jar `WC.jar /user/osboxes/input/computer.txt /user/osboxes/output2/`

Output:-

```

2021-01-12 18:18:59,089 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, re
1      19
10     12
100    22
101    20
102    11
103    12
104     9
105    37
106    36
107    17
108    23
109    38
11     19
110    12
111    14
112    27
113    29
114    50
115    12
116    10
117    10
118    14
119     8

```

4. Write a MapReduce program to find out the dates with maximum temperature greater than 40 (A Hot Day) and minimum temperature lower than 10 (A Cold Day).

Take WeatherData.txt as input file.

Mapper class :

```
HCDMapper.java
1 import java.io.IOException;
2 import org.apache.hadoop.io.LongWritable;
3 import org.apache.hadoop.io.Text;
4 import org.apache.hadoop.mapreduce.Mapper;
5
6 public class HCDMapper extends Mapper<LongWritable, Text, Text, Text>{
7
8     public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException{
9
10         String line = value.toString();
11
12         if(line.length() != 0) {
13
14             String date = line.substring(6,14);
15
16             float temp_max = Float.parseFloat(line.substring(41,47).trim());
17
18             float temp_min = Float.parseFloat(line.substring(48,55).trim());
19
20             if(temp_max > 30.0) {
21
22                 context.write(new Text(date), new Text("Hot Day "+temp_max));
23             }
24
25             if(temp_min < 10.0) {
26                 context.write(new Text(date), new Text("Cold Day "+temp_min));
27             }
28
29         }
30     }
31 }
32
33
34 }
```

Reducer class :

```
HCDReducer.java
1
2 import java.io.IOException;
3 import java.util.Iterator;
4
5 import org.apache.hadoop.io.Text;
6 import org.apache.hadoop.mapreduce.Reducer;
7
8 public class HCDReducer extends Reducer<Text, Text, Text, Text>{
9
10     public void reduce(Text key, Iterator<Text> values, Context context) throws IOException, InterruptedException{
11
12         context.write(key, new Text(values.next().toString()));
13     }
14 }
15
16 }
```

Main class :

```

1 import org.apache.hadoop.io.Text;
2 import org.apache.hadoop.conf.Configuration;
3 import org.apache.hadoop.fs.Path;
4 import org.apache.hadoop.mapreduce.Job;
5 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
6 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
7
8 public class HotColdDay {
9
10     public static void main(String[] args) throws Exception {
11
12         Configuration conf = new Configuration();
13
14         Job job = Job.getInstance(conf, "temprature");
15
16         job.setJarByClass(HotColdDay.class);
17         job.setMapOutputKeyClass(Text.class);
18         job.setMapOutputValueClass(Text.class);
19         job.setMapperClass(HCDMapper.class);
20         job.setReducerClass(HCDReducer.class);
21         job.setOutputKeyClass(Text.class);
22         job.setOutputValueClass(Text.class);
23
24         FileInputFormat.addInputPath(job, new Path(args[0]));
25         FileOutputFormat.setOutputPath(job, new Path(args[1]));
26
27         System.exit(job.waitForCompletion(true) ? 0 : 1);
28
29     }
30
31 }
32
33 }

```

Output :-

```

20200101 Cold Day -1.5
20200102 Cold Day -1.7
20200103 Cold Day 5.4
20200104 Cold Day 5.2
20200105 Cold Day -1.8
20200106 Cold Day -4.9
20200107 Cold Day -4.8
20200108 Cold Day -4.5
20200109 Cold Day -7.0
20200110 Cold Day -7.4
20200111 Cold Day 8.3
20200112 Cold Day 3.9
20200113 Cold Day 1.1
20200114 Cold Day 1.2
20200115 Cold Day -3.1
20200116 Cold Day -2.5
20200117 Cold Day -9.6
20200118 Cold Day -11.7
20200119 Cold Day -2.6
20200120 Cold Day -8.7
20200121 Cold Day -11.0
20200122 Cold Day -12.6
20200123 Cold Day -10.2
20200124 Cold Day -4.5
20200125 Cold Day 1.0
20200126 Cold Day -0.2
20200127 Cold Day -1.9
20200128 Cold Day 1.6
20200129 Cold Day -6.7
20200130 Cold Day -6.9
20200131 Cold Day -7.8
20200201 Cold Day -1.7
20200202 Cold Day -0.3
20200203 Cold Day -4.0

```